

Cold Review & Context Bias

Why a fresh agent catches what a long session can't — and where the drift actually lives.

A long, friendly working session accumulates a quiet pull toward agreement. Cold review is the countermeasure: not a smarter reviewer, a *differently-positioned* one.

1 WHY COLD REVIEW WORKS

A cold reviewer sees only what's **on the page** — not what the author meant. It re-derives intent from the artifact plus the requirements, so it never inherits the author's premises or blind spots.

A requirement misread at step one becomes the author's *premise* — and premises don't get re-examined, they get built on. The reviewer, reconstructing from text alone, lands on exactly the gap the author glided over: the untestable criterion, the assumption never written down.

2 HOW A SESSION DRIFTS

Three pressures grow with the length of a cooperative transcript:

- **Agreement gradient.** A model continues the pattern it's in. A long, successful, friendly exchange makes the *agreeable* next token higher-probability. It strengthens as the history of approval grows.

- **Self-consistency.** Positions taken earlier resist contradiction. Local coherence — not looking like you changed your mind — competes with being right.
- **Ownership.** Jointly-built artifacts acquire a mild "these are good, they're ours" investment a stranger wouldn't feel.

3 THE PLAYBOOK

- **Execute in-session** against a clear spec — the tests are the check, bias is near-zero.
- **Start fresh** (or spawn a cold subagent) for a *new evaluative call*: "is this even the right design?"
- **Hand off files, not chat.** The next stage should inherit the artifact, not the conversational momentum.
- **Unstick in place:** "argue the other side," "review as if you'd never seen it," or route the decision to a reviewer that sees the artifact, not the rapport.

THE CALIBRATION THAT MATTERS

Bias is strongest **exactly where it's hardest to see**: open-ended judgment with no external anchor ("is this good?"). It is **near-zero on anything verifiable** — a test passes or it doesn't; ground truth ignores rapport. *The tell*: resistance is easy when there's a rule or artifact to check against, and hardest on pure judgment late in a warm session.

Honest caveat. A model can't reliably introspect its own bias — don't wait for it to *feel* biased. Assume the bias exists, that it's largest on subjective calls late in a friendly session, and design the countermeasure in. That is the whole point of a brainstorm → spec → implement flow with a cold reviewer at the seam: each handoff strips the gradient and forces re-derivation from the artifact alone.